

基于 keras 和 python librosa 库的乐器声音分类识别

制作人：张斯然 物理 2052972


### 一、背景介绍

物理声学有关的实验中, 用到了频谱分析, 因此, 就上网查找相关资料, 发现 python 的 librosa 库可以用于相关的音频处理, 可以提取到声音有关的相关特征, 如基频信号, 声音的过零率, 梅尔倒谱系数, 频谱滚降点等一系列频谱特征。

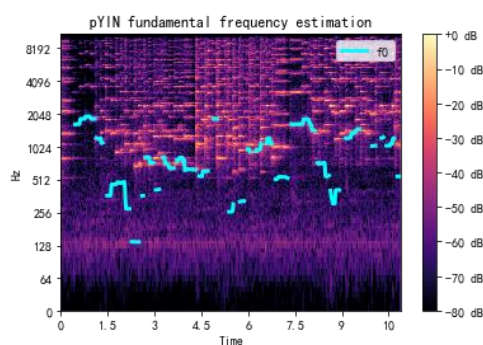
这里以小提琴的一类声音为例。

### 3-1. 小提琴类


#### (1) 波形图




#### (2) 基频分析




#### (3) 短时傅里叶变换得到的时变频谱




#### (4) 过零率



(5) MFCC



(6) 频谱滚降点



## 二、数据集的建立

有了处理声音频谱的方法，只需要有数据即可。于是我自己从网络上下载了钢琴，手风琴，吉他，小提琴，长笛 5 类乐器的独奏曲，将曲目进行切分，每 30s 切分一段，建立了每类大致 25 个片段的数据集。

|             |                  |     |
|-------------|------------------|-----|
| changdi     | 2021/12/16 23:12 | 文件夹 |
| guitar      | 2021/12/15 23:15 | 文件夹 |
| piano       | 2021/12/15 23:15 | 文件夹 |
| shoufengqin | 2021/12/16 23:12 | 文件夹 |
| test        | 2021/12/17 0:00  | 文件夹 |
| violin      | 2021/12/15 23:15 | 文件夹 |

5 类，每类下有 20 余首 30s 的音频曲目。

- 1.wav
- 2.wav
- 3.wav
- 4.wav
- 5.wav
- 6.wav
- 7.wav
- 8.wav
- 9.wav
- 10.wav
- 11.wav
- 13.wav
- 14.wav
- 15.wav
- 16.wav
- 17.wav
- 18.wav
- 19.wav
- 20.wav
- 21.wav
- 22.wav
- 23.wav
- 24.wav
- 25.wav

### 3. 模型搭建与训练

```
In [1]: import librosa
import numpy as np
import os
from keras.models import load_model
from sklearn.preprocessing import StandardScaler
from keras.utils import np_utils
from keras import models
from keras.layers import Dense, Dropout
from keras.models import Sequential # 网络模型
from keras.layers import Dense, Dropout
```

python librosa库有时会出现bug，现实 backend error，此时需要打开ffmpeg包，添加到环境变量，然后在librosa下的aread.py里打开,然后将ffmpeg.py的路径替换。

将训练数据先进行预处理，对每个曲目进行频谱分析，将各个特征的结果保存在array

```
In [2]: classes = 'piano guitar violin changdi shoufengqin'.split()
data_set = []
label_set = []
label2id = {class1:i for i,class1 in enumerate(classes)}
id2label = {i:class1 for i,class1 in enumerate(classes)}
print(label2id)
for c in classes:
    print(c)
    for filename in os.listdir(f'D:/张斯然文件/musicdate/date/classes/{c}/'):
        songname = f'D:/张斯然文件/musicdate/date/classes/{c}/{filename}'
        y, sr = librosa.load(songname, mono=True, duration=30)
        chroma_stft = librosa.feature.chroma_stft(y=y, sr=sr)
        rmse = librosa.feature.rms(y=y)
        spec_cent = librosa.feature.spectral_centroid(y=y, sr=sr)
        spec_bw = librosa.feature.spectral_bandwidth(y=y, sr=sr)
        rolloff = librosa.feature.spectral_rolloff(y=y, sr=sr)
        zcr = librosa.feature.zero_crossing_rate(y)
        mfcc = librosa.feature.mfcc(y=y, sr=sr)

        to_append = f'{np.mean(chroma_stft)} {np.mean(rmse)} {np.mean(spec_cent)} {np
for e in mfcc:
    to_append += f' {np.mean(e)}'
        data_set.append([float(i) for i in to_append.split(" ")])
        label_set.append(label2id[c])
```

```
{'piano': 0, 'guitar': 1, 'violin': 2, 'changdi': 3, 'shoufengqin': 4}
```

```
piano
```

```
C:\Users\cm\anaconda3\lib\site-packages\librosa\core\audio.py:165: UserWarning: PySoundFile failed. Trying audioread instead.
```

```
warnings.warn("PySoundFile failed. Trying audioread instead.")
```

```
C:\Users\cm\anaconda3\lib\site-packages\librosa\core\audio.py:165: UserWarning: PySoundFile failed. Trying audioread instead.
```

```
warnings.warn("PySoundFile failed. Trying audioread instead.")
```

```
guitar
```

```
violin
```

changdi  
shoufengqin

```
In [3]: scaler = StandardScaler()
X = scaler.fit_transform(np.array(data_set, dtype = float))
y = np_utils.to_categorical(np.array(label_set))
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1)
print(X_train.shape[1])
```

26

## 准备模型的建立

```
In [5]: def Model_creat():
        model= Sequential()
        model.add(Dense(units=120,activation='relu',input_shape=(X_train.shape[1],)))#构造
        model.add(Dense(units=60,activation='relu'))
        model.add(Dense(units=30, activation='relu'))
        model.add(Dropout(0.5))
        model.add(Dense(units=5,activation='softmax'))#输出节点
        return model
```

```
In [6]: model=Model_creat()
```

```
In [7]: model.compile(optimizer='adam',loss='categorical_crossentropy',metrics=['accuracy'],'c
```

模型的训练过程还是出现过许多问题，有时候在最后一层使用 **relu**，使得**loss**不能成功计算，通过不停的调整参数和进行比较，得到比较好的结果有**model13,modelshow**,均已经保存在**savemodel**里了

```
In [8]: model.fit(X_train, y_train, epochs=50, batch_size=64)
```

```
Epoch 1/50
2/2 [=====] - 17s 3ms/step - loss: 1.7172 - accuracy: 0.1721
- categorical_accuracy: 0.1721
Epoch 2/50
2/2 [=====] - 0s 5ms/step - loss: 1.5585 - accuracy: 0.2248 -
categorical_accuracy: 0.2248
Epoch 3/50
2/2 [=====] - 0s 4ms/step - loss: 1.4352 - accuracy: 0.2900 -
categorical_accuracy: 0.2900
Epoch 4/50
2/2 [=====] - 0s 2ms/step - loss: 1.3507 - accuracy: 0.3613 -
categorical_accuracy: 0.3613
Epoch 5/50
2/2 [=====] - 0s 2ms/step - loss: 1.2508 - accuracy: 0.4021 -
categorical_accuracy: 0.4021
Epoch 6/50
2/2 [=====] - 0s 2ms/step - loss: 1.1750 - accuracy: 0.5319 -
categorical_accuracy: 0.5319
Epoch 7/50
2/2 [=====] - 0s 2ms/step - loss: 1.1046 - accuracy: 0.6031 -
categorical_accuracy: 0.6031
Epoch 8/50
2/2 [=====] - 0s 2ms/step - loss: 1.1132 - accuracy: 0.6231 -
categorical_accuracy: 0.6231
```

Epoch 9/50  
2/2 [=====] - 0s 2ms/step - loss: 1.0094 - accuracy: 0.7410 -  
categorical\_accuracy: 0.7410  
Epoch 10/50  
2/2 [=====] - 0s 2ms/step - loss: 0.9927 - accuracy: 0.6750 -  
categorical\_accuracy: 0.6750  
Epoch 11/50  
2/2 [=====] - 0s 3ms/step - loss: 0.9153 - accuracy: 0.7500 -  
categorical\_accuracy: 0.7500  
Epoch 12/50  
2/2 [=====] - 0s 2ms/step - loss: 0.8732 - accuracy: 0.7315 -  
categorical\_accuracy: 0.7315  
Epoch 13/50  
2/2 [=====] - 0s 2ms/step - loss: 0.7834 - accuracy: 0.7581 -  
categorical\_accuracy: 0.7581  
Epoch 14/50  
2/2 [=====] - 0s 3ms/step - loss: 0.8340 - accuracy: 0.6194 -  
categorical\_accuracy: 0.6194  
Epoch 15/50  
2/2 [=====] - 0s 2ms/step - loss: 0.7833 - accuracy: 0.6654 -  
categorical\_accuracy: 0.6654  
Epoch 16/50  
2/2 [=====] - 0s 4ms/step - loss: 0.6698 - accuracy: 0.7871 -  
categorical\_accuracy: 0.7871  
Epoch 17/50  
2/2 [=====] - 0s 2ms/step - loss: 0.5482 - accuracy: 0.8635 -  
categorical\_accuracy: 0.8635  
Epoch 18/50  
2/2 [=====] - 0s 2ms/step - loss: 0.5670 - accuracy: 0.7990 -  
categorical\_accuracy: 0.7990  
Epoch 19/50  
2/2 [=====] - 0s 3ms/step - loss: 0.5580 - accuracy: 0.8227 -  
categorical\_accuracy: 0.8227  
Epoch 20/50  
2/2 [=====] - 0s 2ms/step - loss: 0.5233 - accuracy: 0.8531 -  
categorical\_accuracy: 0.8531  
Epoch 21/50  
2/2 [=====] - 0s 2ms/step - loss: 0.4555 - accuracy: 0.8398 -  
categorical\_accuracy: 0.8398  
Epoch 22/50  
2/2 [=====] - 0s 2ms/step - loss: 0.5197 - accuracy: 0.8294 -  
categorical\_accuracy: 0.8294  
Epoch 23/50  
2/2 [=====] - 0s 2ms/step - loss: 0.5246 - accuracy: 0.7990 -  
categorical\_accuracy: 0.7990  
Epoch 24/50  
2/2 [=====] - 0s 2ms/step - loss: 0.4198 - accuracy: 0.8583 -  
categorical\_accuracy: 0.8583  
Epoch 25/50  
2/2 [=====] - 0s 3ms/step - loss: 0.3621 - accuracy: 0.8917 -  
categorical\_accuracy: 0.8917  
Epoch 26/50  
2/2 [=====] - 0s 2ms/step - loss: 0.3034 - accuracy: 0.9140 -  
categorical\_accuracy: 0.9140  
Epoch 27/50  
2/2 [=====] - 0s 2ms/step - loss: 0.3397 - accuracy: 0.8798 -  
categorical\_accuracy: 0.8798  
Epoch 28/50  
2/2 [=====] - 0s 3ms/step - loss: 0.2810 - accuracy: 0.9392 -  
categorical\_accuracy: 0.9392  
Epoch 29/50  
2/2 [=====] - 0s 2ms/step - loss: 0.3446 - accuracy: 0.9035 -  
categorical\_accuracy: 0.9035  
Epoch 30/50  
2/2 [=====] - 0s 2ms/step - loss: 0.2491 - accuracy: 0.9273 -  
categorical\_accuracy: 0.9273  
Epoch 31/50  
2/2 [=====] - 0s 2ms/step - loss: 0.2080 - accuracy: 0.9458 -  
categorical\_accuracy: 0.9458

```

Epoch 32/50
2/2 [=====] - 0s 2ms/step - loss: 0.1663 - accuracy: 0.9681 -
categorical_accuracy: 0.9681
Epoch 33/50
2/2 [=====] - 0s 2ms/step - loss: 0.3205 - accuracy: 0.8494 -
categorical_accuracy: 0.8494
Epoch 34/50
2/2 [=====] - 0s 2ms/step - loss: 0.2212 - accuracy: 0.9458 -
categorical_accuracy: 0.9458
Epoch 35/50
2/2 [=====] - 0s 2ms/step - loss: 0.2342 - accuracy: 0.9287 -
categorical_accuracy: 0.9287
Epoch 36/50
2/2 [=====] - 0s 2ms/step - loss: 0.1529 - accuracy: 0.9681 -
categorical_accuracy: 0.9681
Epoch 37/50
2/2 [=====] - 0s 2ms/step - loss: 0.2097 - accuracy: 0.9310 -
categorical_accuracy: 0.9310
Epoch 38/50
2/2 [=====] - 0s 2ms/step - loss: 0.2377 - accuracy: 0.9392 -
categorical_accuracy: 0.9392
Epoch 39/50
2/2 [=====] - 0s 3ms/step - loss: 0.1451 - accuracy: 0.9615 -
categorical_accuracy: 0.9615
Epoch 40/50
2/2 [=====] - 0s 3ms/step - loss: 0.1643 - accuracy: 0.9681 -
categorical_accuracy: 0.9681
Epoch 41/50
2/2 [=====] - 0s 2ms/step - loss: 0.1477 - accuracy: 0.9577 -
categorical_accuracy: 0.9577
Epoch 42/50
2/2 [=====] - 0s 2ms/step - loss: 0.1530 - accuracy: 0.9392 -
categorical_accuracy: 0.9392
Epoch 43/50
2/2 [=====] - 0s 2ms/step - loss: 0.1980 - accuracy: 0.9206 -
categorical_accuracy: 0.9206
Epoch 44/50
2/2 [=====] - 0s 2ms/step - loss: 0.1347 - accuracy: 0.9444 -
categorical_accuracy: 0.9444
Epoch 45/50
2/2 [=====] - 0s 2ms/step - loss: 0.1469 - accuracy: 0.9644 -
categorical_accuracy: 0.9644
Epoch 46/50
2/2 [=====] - 0s 3ms/step - loss: 0.0998 - accuracy: 0.9815 -
categorical_accuracy: 0.9815
Epoch 47/50
2/2 [=====] - 0s 999us/step - loss: 0.0923 - accuracy: 0.9696 -
categorical_accuracy: 0.9696
Epoch 48/50
2/2 [=====] - 0s 2ms/step - loss: 0.1480 - accuracy: 0.9340 -
categorical_accuracy: 0.9340
Epoch 49/50
2/2 [=====] - 0s 3ms/step - loss: 0.1056 - accuracy: 0.9562 -
categorical_accuracy: 0.9562
Epoch 50/50
2/2 [=====] - 0s 2ms/step - loss: 0.1218 - accuracy: 0.9644 -
categorical_accuracy: 0.9644

```

Out[8]: <keras.callbacks.History at 0x2761b1287c0>

## 保存训练好的模型

In [1]: `#model.save(f'D:/张斯然文件/musicdate/date/modelshow.h5')`

尝试进行验证，选了训练集中没有的5个曲目，保存在数据集目录下classes下的test类里，结果发现还挺准的，但是看概率的话，有一两个可能有点模糊

c——长笛

g——吉他

s——手风琴

p——钢琴

v——小提琴

In [9]:

```
import librosa
import numpy as np
import os
from keras.models import load_model
from sklearn.preprocessing import StandardScaler
from keras.utils import np_utils
model_path = f'D:/张斯然文件/musicdate/date/model13.h5'
# 载入模型
model = load_model(model_path)
data_set1=[]
for filename in os.listdir(f'D:/张斯然文件/musicdate/date/classes/test/'):
    songname = f'D:/张斯然文件/musicdate/date/classes/test/{filename}'
    print(songname)
    y, sr = librosa.load(songname, mono=True, duration=30)
    chroma_stft = librosa.feature.chroma_stft(y=y, sr=sr)
    rmse = librosa.feature.rms(y=y)
    spec_cent = librosa.feature.spectral_centroid(y=y, sr=sr)
    spec_bw = librosa.feature.spectral_bandwidth(y=y, sr=sr)
    rolloff = librosa.feature.spectral_rolloff(y=y, sr=sr)
    zcr = librosa.feature.zero_crossing_rate(y)
    mfcc = librosa.feature.mfcc(y=y, sr=sr)

    to_append = f'{np.mean(chroma_stft)} {np.mean(rmse)} {np.mean(spec_cent)} {np.me

    for e in mfcc:
        to_append += f' {np.mean(e)}'

    data_set1.append([float(i) for i in to_append.split(" ")])

scaler = StandardScaler()
X1= scaler.fit_transform(np.array(data_set1, dtype = float))
result=model.predict(X1)
array=np.argmax(result,axis=1)
print(result)
print(array)
classes=['钢琴','吉他','小提琴','长笛','手风琴']
for j in range(0,len(array)):
    print("第{}个为".format(j+1),classes[array[j]])
```

D:/张斯然文件/musicdate/date/classes/test/c1.mp3

C:\Users\cm\anaconda3\lib\site-packages\librosa\core\audio.py:165: UserWarning: PySoundFile failed. Trying audioread instead.

warnings.warn("PySoundFile failed. Trying audioread instead.")



```
D:/张斯然文件/musicdate/date/classes/test/gl.mp3
D:/张斯然文件/musicdate/date/classes/test/pl.mp3
D:/张斯然文件/musicdate/date/classes/test/sl.mp3
D:/张斯然文件/musicdate/date/classes/test/vl.mp3
[[1.84158352e-03 1.42207171e-03 1.31978682e-04 9.96122897e-01
  4.81391791e-04]
 [2.12312430e-01 7.36021459e-01 5.59710013e-03 1.47150655e-03
  4.45974581e-02]
 [8.45506310e-01 1.22267537e-01 5.69243683e-03 1.21670775e-02
  1.43666435e-02]
 [3.53744836e-03 2.76785381e-02 3.68632287e-01 1.07908566e-02
  5.89360833e-01]
 [3.21190059e-02 6.64945394e-02 6.67403102e-01 4.00314108e-03
  2.29980201e-01]]
[3 1 0 4 2]
第1个为 长笛
第2个为 吉他
第3个为 钢琴
第4个为 手风琴
第5个为 小提琴
```

In [ ]: